

Programming Assignment - 2:

MLFQ Mutex Implementation

Mehmet Altunören

22539

1. Overview

In this report, I describe two key components I developed for PA2:

1. A **thread-safe concurrent queue** (queue.h), and
2. An **MLFQ-based mutex** (MLFQmutex.h).

I explain my design choices, demonstrate how each implementation satisfies the correctness requirements under concurrency, and argue how they achieve fairness and performance goals.

2. Concurrent Queue Implementation (queue.h)

2.1 Algorithm Choice

To implement a thread-safe FIFO queue, I adopted the classic **two-lock linked-list queue** pattern (Michael & Scott style). I chose this because:

- It cleanly separates enqueue and dequeue work under distinct locks.
- It avoids heavy contention: enqueueers lock only the tail, while dequeueers lock only the head.
- It is simple to reason about and adapt in C++ with `pthread_mutex_t`.

2.2 Internal Data Structures

- **Node<T>**: each node stores a T value and a Node* next.
- **Dummy node**: on construction, I allocate one dummy node that both head and tail point to.
- **Pointers**:
 - head always refers to the dummy or the node preceding the first real element.
 - tail refers to the last real node (or the dummy if empty).
- **Mutexes**:
 - head_lock (a pthread_mutex_t) protects operations on the head pointer and its next link.
 - tail_lock protects operations on the tail pointer and its next link.

2.3 API Methods

Constructor & Destructor

- **Constructor**: allocates the dummy node, initializes both locks.
- **Destructor**: iterates from head, deleting each node; destroys both mutexes.

enqueue(const T& item)

1. Allocate a new Node<T> containing item.

2. Lock `tail_lock`.
3. Link the new node at `tail->next` and set `tail = newNode`.
4. Unlock `tail_lock`.

Because only enqueueers ever touch `tail`, multiple threads can call `enqueue()` concurrently without interfering with dequeuers.

`dequeue()` → T

1. Lock `head_lock`.
2. Let `first = head->next`.
 - If `first == nullptr`, the queue is empty: unlock `head_lock` and throw `std::runtime_error("Queue is empty!")`.
3. Copy `first->value` to a local variable.
4. Advance `head` to `first`, delete the old dummy node.
5. Unlock `head_lock`.
6. Return the copied value.

Only one thread at a time can manipulate `head`, ensuring mutual exclusion for removals.

`isEmpty()` / `print()`

- **`isEmpty()`** locks `head_lock`, checks `head->next == nullptr`, unlocks, and returns the result.
- **`print()`** locks `head_lock` and iterates from `head->next`, printing each element (or “Empty” if none), then unlocks.

2.4 Thread-Safety and FIFO Guarantee

- By partitioning access, enqueues and dequeues proceed entirely in parallel when the queue is non-empty and not changing `head` and `tail` simultaneously.
- No element can be lost or duplicated: each node is either safely linked under `tail_lock`, or unlinked under `head_lock`.
- FIFO order is enforced because nodes are always appended at the tail and removed from the head in arrival order.

2.5 Performance Considerations

- Lock granularity is minimal: only single operations under each lock.
 - Contention only occurs when many threads enqueue at once (contending for `tail_lock`) or dequeue at once (contending for `head_lock`).
 - Average enqueue/dequeue latency is low and scales well under moderate concurrency.
-

3. MLFQ-Based Mutex Implementation (MLFQmutex.h)

3.1 MLFQ Algorithm Recap

I embedded a **user-level Multi-Level Feedback Queue** scheduler into a pthread mutex to ensure fairness among waiting threads:

1. **Initial priority:** each thread's level (`threadPriority`) is 0 on its first `lock()`.
2. **Quantum:** a fixed time slice Q (in seconds) specified at constructor time.
3. **Priority update:** after exiting the critical section, measure the elapsed time `execTime`; set $\text{newLevel} = \min(\text{oldLevel} + \lfloor \text{execTime} / Q \rfloor, \text{levels} - 1)$.
4. **Selection:** on `unlock()`, pick the **lowest non-empty** priority queue (round robin within the level) to `unpark` next.

3.2 Data Structures

- **thread_local int threadPriority:** tracks each thread's current level across invocations.
- **atomic_flag mutexFlag:** the core lock bit—`test_and_set(memory_order_acquire)` acquires, `clear(memory_order_release)` releases.
- **SpinLock internalLock:** a simple `atomic_flag`-based spin-lock guarding all shared state (queues, bitmask, flag).
- **vector<Queue<pthread_t*>> priorityQueues:** one queue per level to hold parked threads.
- **unsigned activeQueuesMask:** a bitmask where bit $i = 1$ if `priorityQueues[i]` is non-empty; used with `__builtin_ctz` to find the lowest non-empty level in $O(1)$.
- **Garage garage:** provides `setPark()`, `park()`, and `unpark(pthread_t)` to block/wake threads without busy-waiting.
- **Timing fields:** `startTime` and `endTime` record the critical-section entry/exit times.

3.3 lock() Design

```
void lock() {
    internalLock.lock();
    if (!mutexFlag.test_and_set(memory_order_acquire)) {
        // Fast path: acquired immediately
        internalLock.unlock();
    } else {
        // Slow path: someone else holds the lock
        enqueueThread(); // enqueues under internalLock, prints & flushes
        garage.setPark(); // prevent missed wakeup
        internalLock.unlock();
        garage.park();    // block until unparked
    }
    // Record entry time for duration accounting
    startTime = high_resolution_clock::now();
}
```

```
}
```

- **Fast path:** zero-overhead acquired when uncontended.
- **Slow path:** under the guard lock, enqueue self at our current threadPriority, call setPark(), then unlock and park(). This sequence avoids the “lost wake-up” problem.

Printing: enqueueThread() writes “Adding thread with ID: <tid> to level <level>”, and flushes immediately to avoid interleaved console output.

3.4 unlock() Design

```
void unlock() {
    internalLock.lock();
    endTime = high_resolution_clock::now();
    auto duration = duration_cast<seconds>(endTime - startTime);
    adjustThreadPriority(duration);    // bump level by floor(duration/Q)
    pthread_t next = dequeueNextThread(); // finds lowest non-empty queue
    if (next != (pthread_t)-1) {
        garage.unpark(next);        // hand off to a waiting thread
    } else {
        mutexFlag.clear(memory_order_release); // no waiters: fully release
    }
    internalLock.unlock();
}
```

- Compute actual critical-section length, adjust the calling thread’s threadPriority.
- dequeueNextThread() uses activeQueuesMask and __builtin_ctz to locate the next thread in minimal time.

3.5 Correctness

- **Mutual exclusion:** only one thread can clear test_and_set and proceed unparked; internalLock serializes all shared-state mutations.
- **Deadlock-freedom:** every unlock() either unparks a waiter or clears the flag, ensuring progress.

3.6 Fairness

- Threads that run short critical sections (I/O-like) remain at high priority (level 0), minimizing their wait time.
- CPU-bound threads sink to lower levels by $\lfloor \text{execTime}/Q \rfloor$ each time, letting short tasks “overtake” them.
- No thread can starve indefinitely because a long CPU-bound thread cannot stay at the top forever.

3.7 Performance

- **Fast path:** `test_and_set` with no extra locking when uncontended.
 - **Thresholded spin-lock:** `SpinLock` yields to the scheduler after 1000 spins to avoid wasteful busy-waiting.
 - **Park/unpark:** threads block in the OS rather than spin, eliminating CPU overhead under contention.
-

4. Use of Provided `park.h`

I utilized the supplied `Garage` implementation without modification. It correctly handles `park()`, `setPark()`, and `unpark()` to avoid lost wake-ups.

5. Conclusion

- My **concurrent queue** preserves strict FIFO under heavy concurrency while keeping enqueue/dequeue paths separate and low-overhead.
- My **MLFQ mutex** combines a fast, optimistic test-and-set path with a user-level feedback queue scheduler to balance correctness, fairness, and performance.